



Kubernetes

Observability Challenge Solutions

2.1. CHALLENGE: kubectl top

Using the `kubectl top` command and its help text available from `-h` or `--help`, try the following operations:

- Retrieve the current CPU and Memory utilization of all control plane components in the `kube-system` namespace by their label

Discover the control plane labels by getting the pods in the `kube-system` namespace:

```
~$ kubectl -n kube-system get po --show-labels
```

NAME	READY	STATUS	RESTARTS
AGE LABELS			
coredns-76f75df574-9k4tx	1/1	Running	0
101m k8s-app=kube-dns,pod-template-hash=76f75df574			
coredns-76f75df574-9p29r	1/1	Running	0
101m k8s-app=kube-dns,pod-template-hash=76f75df574			
etcd-ip-172-31-4-161	1/1	Running	0
101m component=etcd,tier=control-plane			
kube-apiserver-ip-172-31-4-161	1/1	Running	0
101m component=kube-apiserver,tier=control-plane			
kube-controller-manager-ip-172-31-4-161	1/1	Running	0
101m component=kube-controller-manager,tier=control-plane			
kube-proxy-kbdll	1/1	Running	0
101m controller-revision-hash=5f6677ccc4,k8s-app=kube-proxy,pod-template-generation=1			
kube-scheduler-ip-172-31-4-161	1/1	Running	0
101m component=kube-scheduler,tier=control-plane			
metrics-server-c5fbf4cb9-lr2kd	1/1	Running	0
3m10s k8s-app=metrics-server,pod-template-hash=c5fbf4cb9			
weave-net-x76wg	2/2	Running	1 (98m ago)
98m controller-revision-hash=d94c4d6bc,name=weave-net,pod-template-generation=1			

```
~$
```

The control plane components have a common label: `tier=control-plane`; use it to filter pods

```
~$ kubectl -n kube-system top pods -l tier=control-plane
```

NAME	CPU(cores)	MEMORY(bytes)
etcd-ip-172-31-4-161	21m	42Mi
kube-apiserver-ip-172-31-4-161	47m	281Mi
kube-controller-manager-ip-172-31-4-161	15m	50Mi
kube-scheduler-ip-172-31-4-161	4m	20Mi

```
~$
```

- Sort the output by memory use

```
~$ kubectl -n kube-system top pods -l tier=control-plane --sort-by memory
```

NAME	CPU(cores)	MEMORY(bytes)
kube-apiserver-ip-172-31-4-161	51m	281Mi
kube-controller-manager-ip-172-31-4-161	17m	50Mi
etcd-ip-172-31-4-161	22m	42Mi
kube-scheduler-ip-172-31-4-161	4m	20Mi

```
~$
```

- Which control plane component is currently using the most memory?

Based on the example, the **kube-apiserver** pod. It depends on which one is sorted to the top

- Retrieve a listing of resource use of all nodes in your cluster sorted by CPU utilization without printing the column headers

```
~$ kubectl top nodes --no-headers --sort-by cpu
```

```
ip-172-31-4-161    201m    10%    1782Mi    47%
```

```
~$
```

3.1. CHALLENGE: HPA editing

Before you clean up, try the following operations:

- Adjust the deployment or HPA so that the **web1** deployment scales when a pod reaches 50m of CPU:

Adjusting the deployment and leaving the HPA alone: Change the cpu request in the web1 deployment to **100m**

```
~/hpa$ nano target-deploy.yaml && cat target-deploy.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: web1
  name: web1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: web1
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: web1
    spec:
      containers:
      - image: rxmllc/target
        imagePullPolicy: Never
        name: target
        resources:
          requests:
            cpu: "100m"
```

```
~/hpa$ kubectl apply -f target-deploy.yaml

deployment.apps/web1 configured

~/hpa$
```

OR:

Adjusting the HPA and leaving the deployment alone: Edit the HPA to scale at 25% CPU:

```
~/hpa$ kubectl edit hpa web1
```

```
...

spec:
  maxReplicas: 5
  metrics:
  - resource:
    name: cpu
```

```
    target:
      averageUtilization: 25
      type: Utilization
    type: Resource
status:
conditions:
```

```
:wq
```

```
horizontalpodautoscaler.autoscaling/web1 edited
```

```
~/hpa$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
web1	Deployment/web1	0%/25%	1	5	1	100m

```
~/hpa$
```

- Ensure the HPA can scale the target deployment up to 3 replicas

Edit the HPA so its maxReplicas is 3:

```
~/hpa$ kubectl edit hpa web1
```

```
spec:
  maxReplicas: 3
  metrics:
  - resource:
      name: cpu
      target:
        averageUtilization: 25
        type: Utilization
      type: Resource
  minReplicas: 1
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: web1
```

```
:wq
```

```
horizontalpodautoscaler.autoscaling/web1 edited
```

```
~/hpa$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
web1	Deployment/web1	0%/25%	1	3	1	23m

~/hpa\$

- Adjust the HPA minimum count to 2 replicas:

```
~/hpa$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
driver	1/1	Running	1 (12m ago)	15m
web1-658f9667d6-w4c82	1/1	Running	0	24m

```
~/hpa$ kubectl get deploy web1
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
web1	1/1	1	1	24m

```
~/hpa$ kubectl edit hpa web1
```

```
...

spec:
  maxReplicas: 3
  metrics:
  - resource:
      name: cpu
      target:
        averageUtilization: 25
        type: Utilization
      type: Resource
  minReplicas: 2
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: web1
```

```
:wq

horizontalpodautoscaler.autoscaling/web1 edited

~/hpa$ kubectl get deploy web1 -w
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
web1	1/1	1	1	24m
web1	1/2	1	1	25m

```
web1 1/2 1 1 25m
web1 1/2 1 1 25m
web1 1/2 2 1 25m
web1 2/2 2 2 25m
```

```
^C
```

```
~/hpa$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
driver	1/1	Running	1 (13m ago)	17m
web1-658f9667d6-hjnph	1/1	Running	0	20s
web1-658f9667d6-w4c82	1/1	Running	0	25m

```
~/hpa$
```

- What happens to the deployment?

The deployment's replica count is adjusted up to match the HPA, effectively enforcing one-way scaling.

Copyright (c) 2023-2024 RX-M LLC, Cloud Native & AI Training and Consulting, all rights reserved